

Assignment 3

Restricted Boltzmann Machines on MNIST

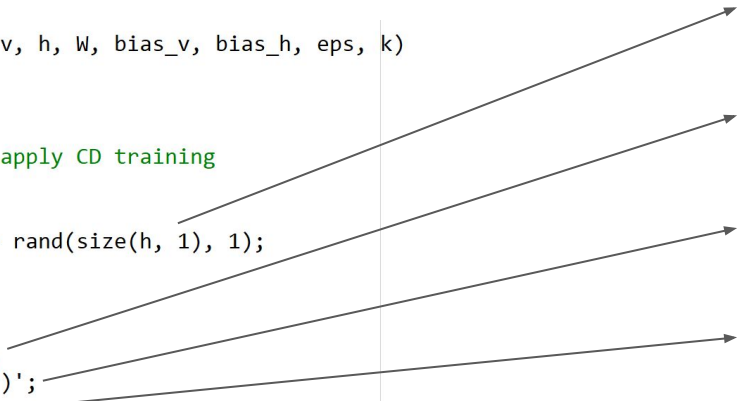
Tarasco Pier Paolo

ID: 619622

Code: <https://github.com/Pier297/ISPR/tree/main/Midterm%202>

Code

```
function [W, bias_v, bias_h] = mini_batch(X, v, h, W, bias_v, bias_h, eps, k)
    delta_W = zeros(size(v, 1), size(h, 1));
    delta_bias_v = zeros(size(v, 1), 1);
    delta_bias_h = zeros(size(h, 1), 1);
    % For each data point of the mini-batch, apply CD training
    for i = 1:size(X, 1)
        v_0 = X(i, :)' ;
        h_0 = logistic(v_0' * W + bias_h)' > rand(size(h, 1), 1);
        % k step gibbs sampling
        h_k = h_0;
        for s = 1:k
            v_k = logistic(W * h_k + bias_v);
            p_j = logistic(v_k' * W + bias_h)' ;
            h_k = p_j > rand(size(h, 1), 1);
        end
        h_k = p_j;
        % When computing the gradient use p_j instead of h_j,
        % this reduces the sampling noise -> leads to faster training
        % -- [Hinton, A Practical Guide to Training Restricted Boltzmann Machines]
        delta_W = delta_W + (v_0 * h_0') - (v_k * h_k');
        delta_bias_v = delta_bias_v + (v_0 - v_k);
        delta_bias_h = delta_bias_h + (h_0 - h_k);
    end
    % Update W
    W = W + (eps / size(X, 1)) * delta_W;
    % Update biases
    bias_v = bias_v + (eps / size(X, 1)) * delta_bias_v;
    bias_h = bias_h + (eps / size(X, 1)) * delta_bias_h;
end
```



$h \sim P(h_j = 1|v)$

$P(v_i = 1|h)$


$P(h_j = 1|v)$

$h \sim P(h_j = 1|v)$

end

Results

100 features of the RBM



```
n_hidden_units = 100;  
k = 1;  
BATCH_SIZE = 20;  
eps = 0.1;  
MAX_EPOCHS_RBM = 20;
```

Softmax Layer

# EPOCHS	TR accuracy	TS accuracy	Time [s]
10	0.921	0.919	240
20	0.923	0.919	450

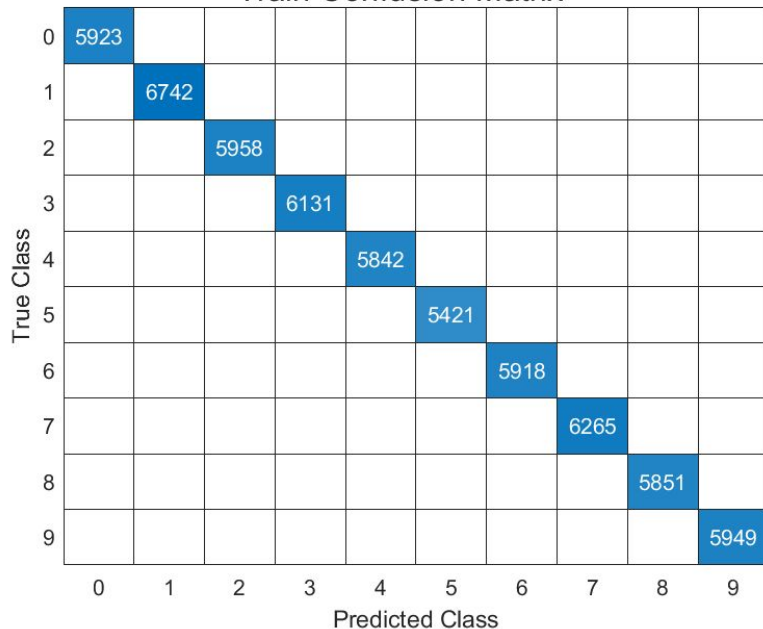
Logistic Regression (with 'fitcnet')

# Hidden Layers	Structure Hidden Layer	TR acc.	TS acc.	Time [s]
0	-	0.811	0.814	35
1	[10]	0.932	0.929	36
1	[100]	1	0.963	36
2	[100, 50]	1	0.964	50
2	[100, 75]	1	0.966	54

Results: `fitcnet(enc_X_train, Y_train, "LayerSizes", [100 75]);`

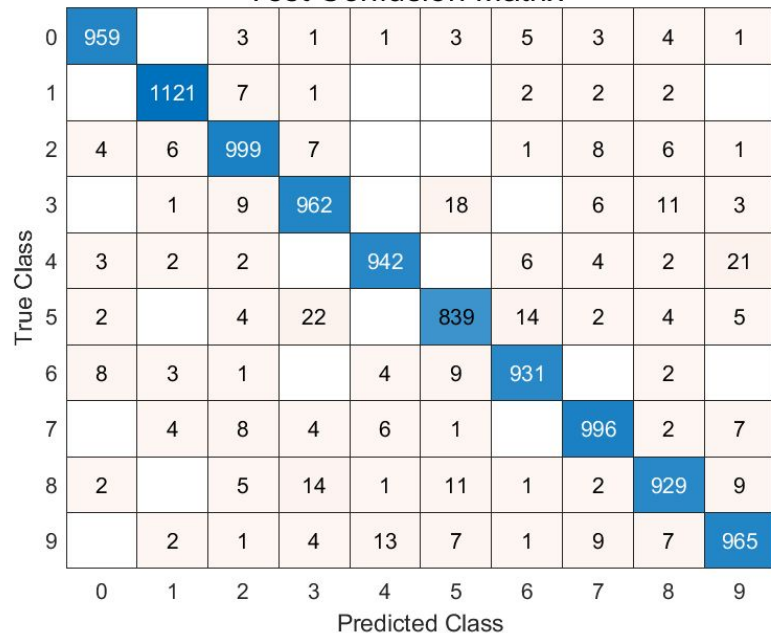
Accuracy = 1

Train Confusion Matrix



Accuracy = 0.966

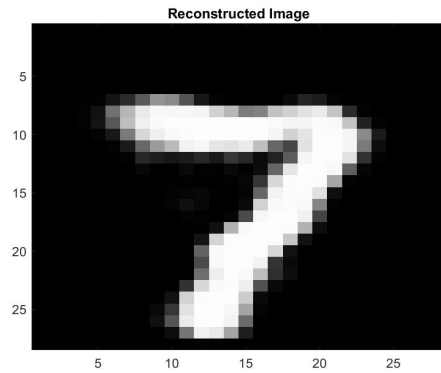


Test Confusion Matrix





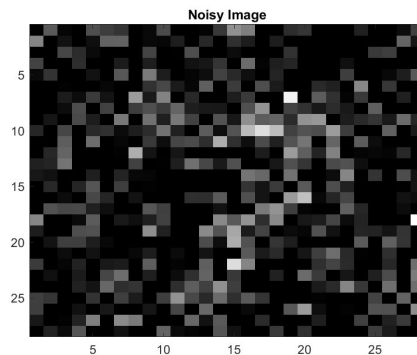
What if we didn't use the RBM representation? **TR acc. = 1** **TS acc. = 0.975**

- **1% performance tradeoff versus 90% compression rate**




Reconstruction



Denoising




Completion




Dreaming


Original Image




Reconstructed Image 1




Reconstructed Image 2




Reconstructed Image 3



Reconstructed Image 4



Energy function



Final considerations:

- Training an RBM takes time and the performance on the classification task decreases, but:
- We can achieve almost the same performance with a 90% compression rate!
- We learn the probability distribution $P(v, h)$